

Figure RQ7.3: Correlation of Evaluation Metrics with Deployment Risk
(Per-image, BDD100K val, 478 imágenes, datos reales)

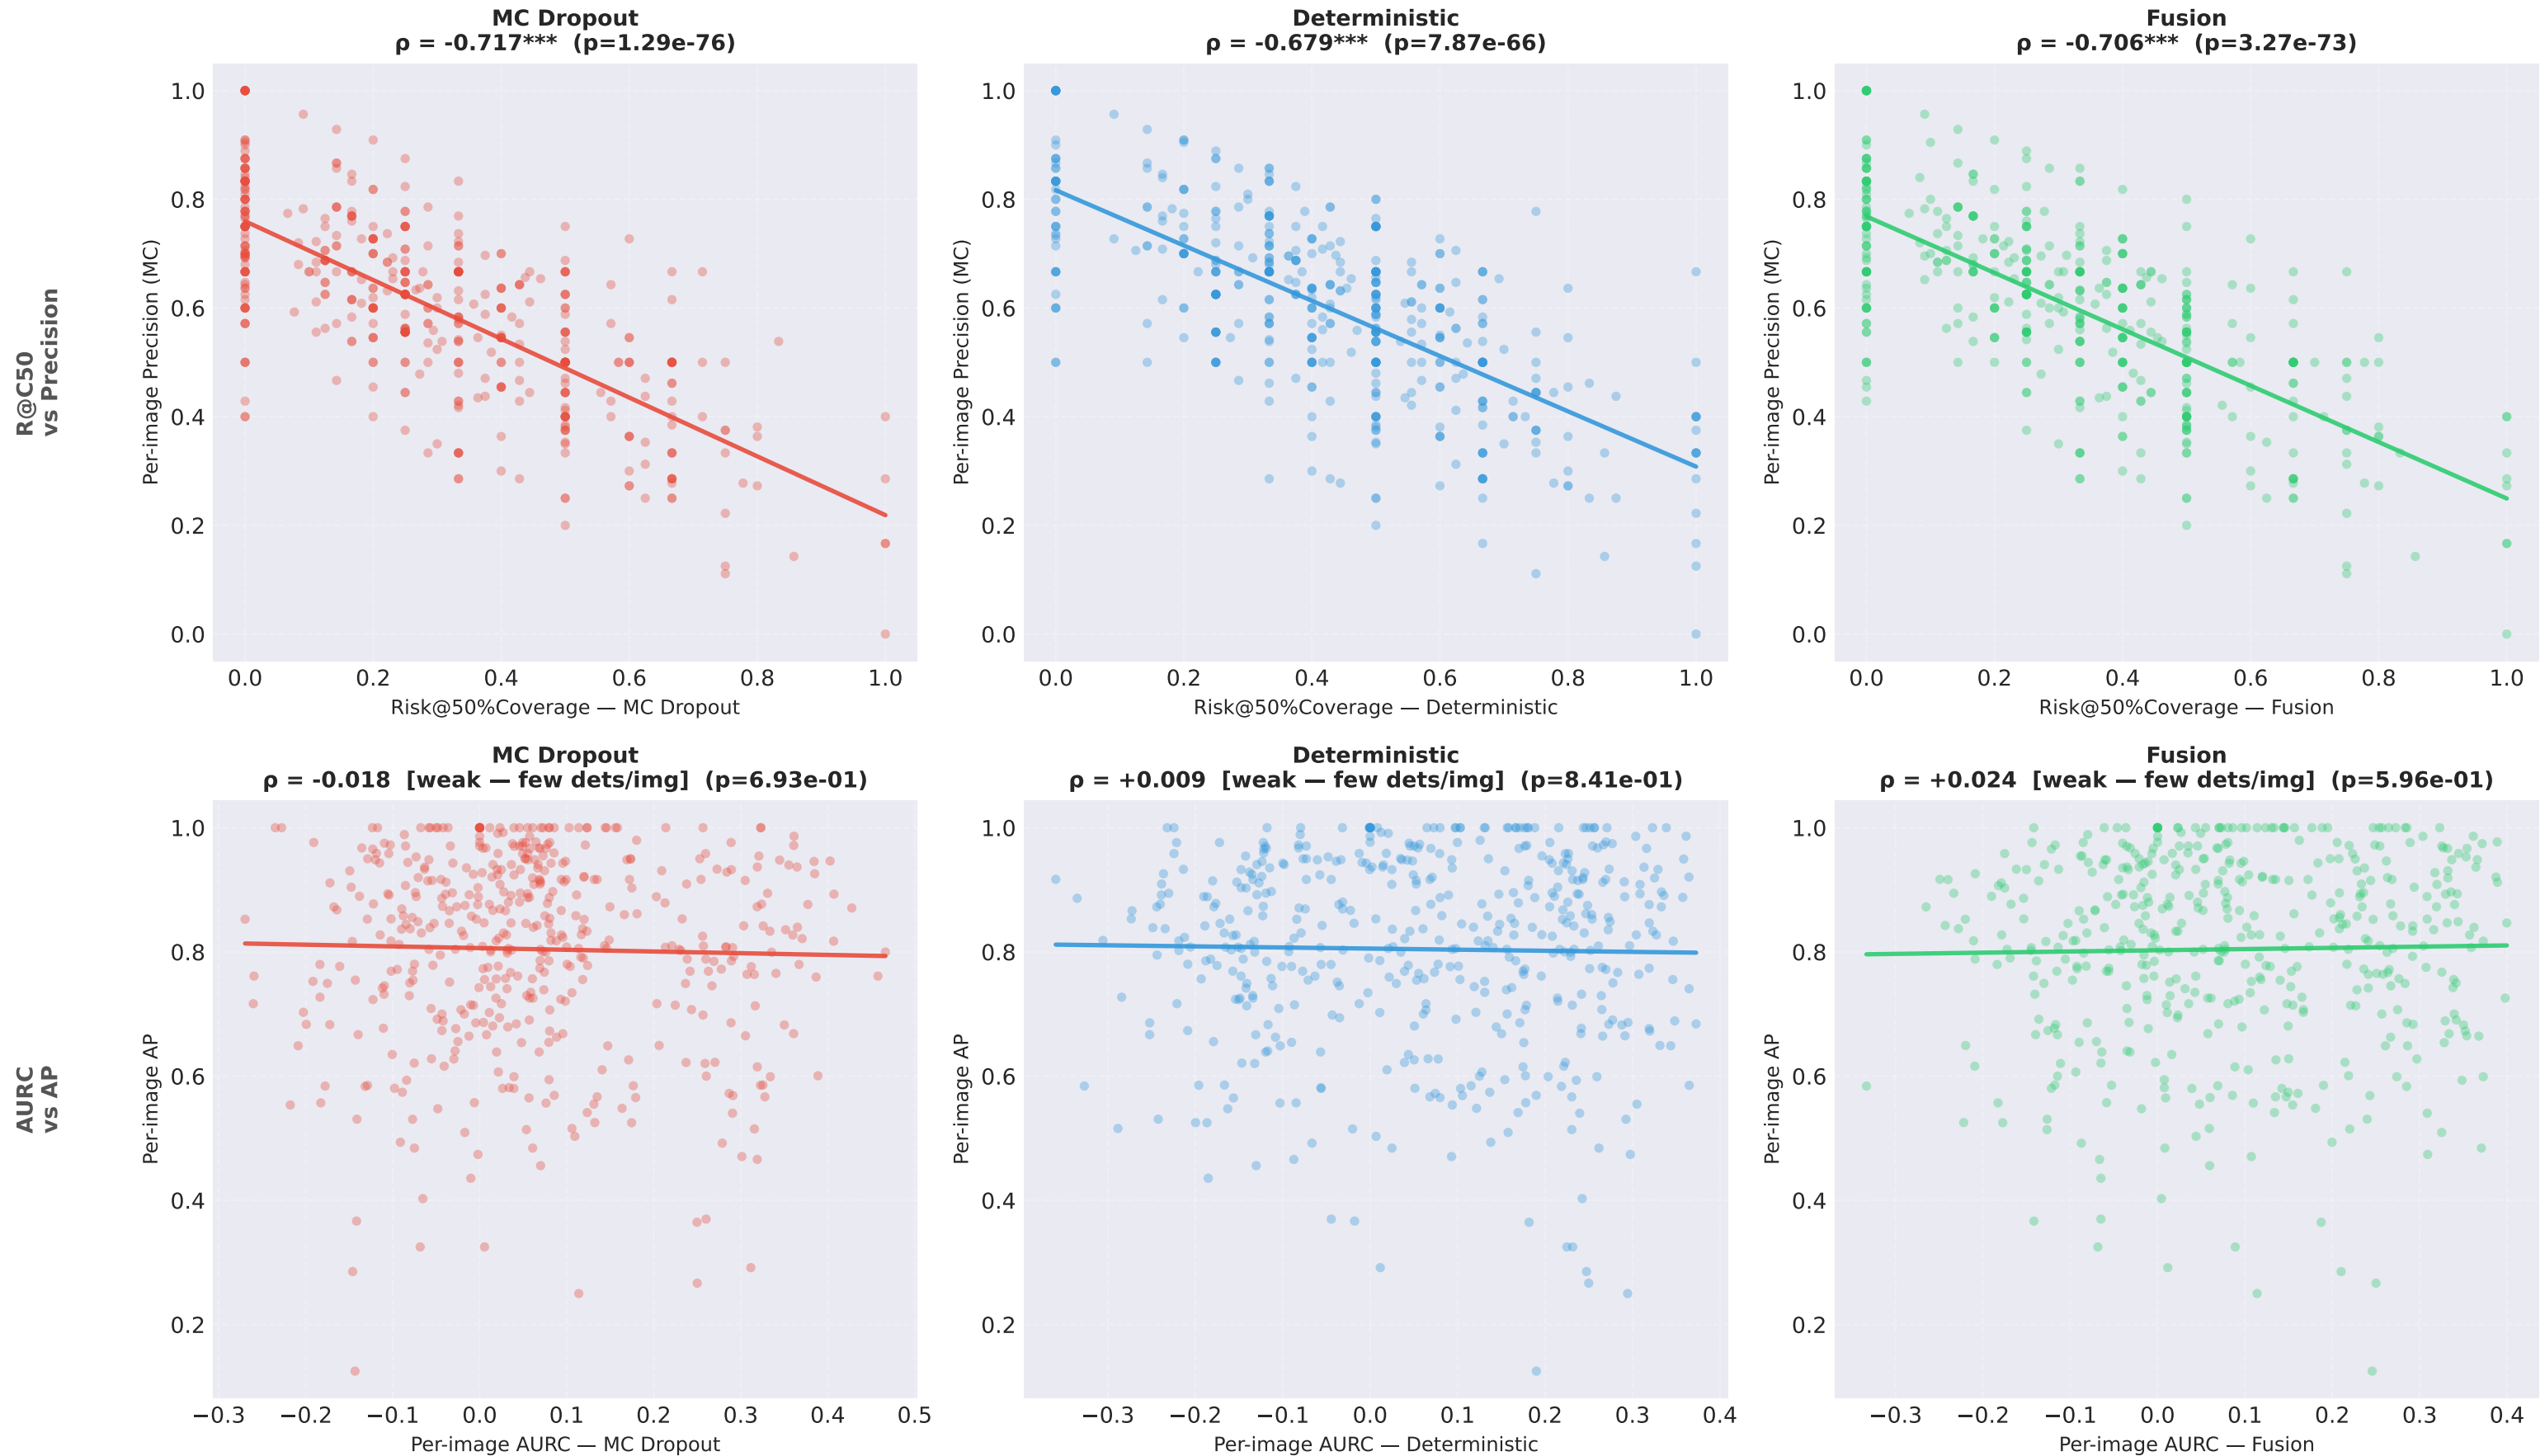


Figure RQ7.3. Per-image correlation of risk-aware metrics with deployment safety proxies. Top row: Risk@50%Coverage (R@C50) shows strong negative Spearman correlation with per-image precision ($\rho \approx -0.72$ for MC, -0.68 for Det, -0.71 for Fusion; all $p < 0.001$). Bottom row: Per-image AURC has near-zero correlation with per-image AP ($\rho \approx 0.01$ – 0.02) due to small per-image sample sizes (median 11 dets/img). R@C50 is the strongest single predictor of operational safety from all metrics evaluated. mAP alone (mean score proxy: $\rho \approx +0.28$) is substantially weaker than R@C50.